

ERNM Vignette

Ian Fellows

November 5, 2012

1 Introduction

A graph is a collection of nodes, each of which may either be connected or not connected to each other node. For the purposes of this document, nodes may not be connected to themselves. In addition to the graph connections, each node may have characteristics which are of interest to the researcher. A network is defined as the union of a graph and the nodal characteristics. Random graphs, where connections between nodes are random but nodal characteristics are either fixed or missing, have a long history in the mathematical literature starting with the simple Erdos-Renyi model (?), and including the more general exponential-family random graph models (ERGM) for which inference requires modern Markov Chain Monte Carlo (MCMC) methods (??). On the other hand we have Gibbs/Markov random field models where nodal attributes are random but interconnections between nodes are fixed. A simple example is the Ising model of ferromagnetism (?) from the statistical physics literature which is exactly solvable under certain network configurations (?); however, most field models require more complex methodologies for inference (?).

In the social network literature, these two classes of models are conceptually defined as “social selection” and “social influence” models. In social selection models, the probability of social ties between individuals are determined by nodal characteristics such as age or sex (see ? and references therein). In social influence models, individuals’ nodal characteristics are determined by social ties (see ? and references therein). ? argues that the processes of tie selection and nodal variate influence are co-occurring phenomena, with ties affecting nodal variates and visa versa, and should therefore be considered together. This chapter presents a joint exponential-family model of connections between nodes (dyads), and nodal attributes, thus representing a unification of social selection and influence. We will refer to this model as an exponential-family random network model (ERNM).

Let the graph Y be an n by n matrix whose entries $Y_{i,j}$ indicate whether subject i and j are connected, where n is the size of the population. Further let X be an $n \times q$ matrix of nodal variates. We define the network to be the random variable (Y, X) .

ERGM:

$$P(Y = y | \eta, X = x) = \frac{1}{c(\eta, x)} e^{\eta \cdot g((y, x)) + o((y, x))}$$

Gibbs/Markov random field:

$$P(X = x | \eta, Y = y) = \frac{1}{c(\eta, y)} e^{\eta \cdot g((y, x)) + o((y, x))}$$

(NEW!!) Exponential-Family Random Network Model (ERNM):

$$P(Y = y, X = x | \eta) = \frac{1}{c(\eta)} e^{\eta \cdot g((y, x)) + o((y, x))}$$

As `ernm` models subsume both `ergm` and Gibbs random field models, the `ernm` package can fit and simulate from both.

```
> set.seed(1)
> library(ernm)
> library(ergm)
```

2 ERGM

In this section we compare an `ergm` fit with the `ergm` package with an `ernm` model fit with the `ernm` package.

```
> data(sampson)
> #fit a model with ergm
> ergmFit <- ergm(samplike ~ edges + asymmetric + match("group"))
```

```
Iteration 1 of at most 20:
Convergence test P-value: 9.7e-03
The log-likelihood improved by 0.001176
Iteration 2 of at most 20:
Convergence test P-value: 1.2e-01
The log-likelihood improved by 0.0005834
Iteration 3 of at most 20:
Convergence test P-value: 5.8e-01
Convergence detected. Stopping.
The log-likelihood improved by 0.0001227
```

This model was fit using MCMC. To examine model diagnostics and check for degeneracy, use the

```
> summary(ergmFit)
```

```
=====
Summary of model fit
=====
```

Formula: `samplike ~ edges + asymmetric + match("group")`

Iterations: 20

Monte Carlo MLE Results:

	Estimate	Std. Error	MCMC %	p-value
edges	-1.5493	0.2296	0	< 1e-04 ***
asymmetric	-0.7162	0.2404	0	0.00312 **
nodematch.group	2.0333	0.3125	0	< 1e-04 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Null Deviance: 424.2 on 306 degrees of freedom

Residual Deviance: 692.5 on 303 degrees of freedom

AIC: 698.5 BIC: 709.7 (Smaller is better.)


```
> #fit a similar model with ernm
> net <- as.BinaryNet(samplike)
> plot(net,vertex.col=net[["group"]],main="Sampson's Monks")
> ernmFit <- ernm(net ~ edges() + reciprocity() + nodeMatch("group"),verbose=0)
```

calculating:

```
> summary(ernmFit)
```

	theta	se	z	p.value
edges	-2.218064	0.2215872	-10.009892	1.379027e-23
reciprocity	1.358206	0.4684402	2.899422	3.738514e-03
nodeMatch	2.000028	0.2995928	6.675820	2.458539e-11

Sampson's Monks



We note a few things here. First, though the model formulas are similar, they are distinct, with `ernm` having its own statistics. The edges and match statistics are equivalent, but the `ernm` model counts symmetric ties, whereas the `ergm` package counts asymmetric ties. Which variates are considered random is specified after the `|` character. by default, only the graph is considered random. If the graph and a variable named "var" is missing, you can use something like:

```
net ~ edges() | var
if only var is random (fixed graph), then you can use noDyad
net ~ edges() | noDyad + var
constraints and offsets are nested in a flag as in:
net ~ edges() + constraint(boundedDegree(1,10))
net ~ edges() + offset(offsetName())
```

Secondly, a new network data structure is used (`BinaryNet`), which can be operated on both from C++ and R. More on this later.

2.1 ERGM Timings

It may be instructive to compare the speed of `ergm` v.s. `ernm`. To do this, we use a model with a term for edges and one for triangles. the parameter for the triangle term is set to 0, and the edges parameter is chosen such that the mean degree of the network is 10. Both packages are set to use Tie-Dyad toggling, with a burn in of 1000000, an interval of 1000, and a sample size of 1000.

```

> library(ggplot2)
> ns <- c(25,50,100,500,1000,2000,3500,5000,7500,10000)
> ergmTimings <- c()
> ernmTimings <- c()
> for(n in ns){
+   print(n)
+   p <- 10/(n-1)
+   theta <- c(log(p/(1-p)),0)
+
+   nw <- network.initialize(n, directed = FALSE)
+   net <- as.BinaryNet(nw)
+
+   sampler <- createCppSampler(net ~ edges() + triangles(),sampler="VertexMetropolis")
+   sampler$getModel()$setThetas(theta)
+
+   ernmTime <- system.time(sims1 <- sampler$generateSampleStatistics(1000000,1000,1000))
+
+   ergmTime <- system.time(sims <- simulate(nw ~ edges + triangle,nsim=1000,coef=theta,stat
+     control = control.simulate.formula(MCMC.burnin = 1000000, MCMC.interval = 1000))
+   ergmTimings <- c(ergmTimings,ergmTime[3])
+   ernmTimings <- c(ernmTimings,ernmTime[3])
+ }

[1] 25
[1] 50
[1] 100
[1] 500
[1] 1000
[1] 2000
[1] 3500
[1] 5000
[1] 7500
[1] 10000

> print(ergmTimings)


elapsed elapsed elapsed elapsed elapsed elapsed elapsed elapsed elapsed elapsed
 1.635   1.772   1.861   2.503   3.191   4.747   6.984  13.832  19.164  23.882

> print(ernmTimings)

elapsed elapsed elapsed elapsed elapsed elapsed elapsed elapsed elapsed elapsed
 2.681   2.426   2.303   2.270   2.652   2.572   2.804   2.955   3.571   4.226

> dat <- rbind(data.frame(time=ergmTimings,n=ns,package="ergm"),
+               data.frame(time=ernmTimings,n=ns,package="ernm"))
> p <- qplot(x=n,y=time,data=dat,color=package,geom="line")+geom_point()+ theme_bw()
> print(p)

```



Note how the ergm simulation times explode up to around 30 as the number of nodes increases, whereas the ernm simulations remain under 5 seconds. This is because all of ernm's toggling / change statistics are constant time as n increases provided the average degree remains constant. These results represent a marked improvement for ergm, previous versions of ergm that I tested yielded results over 2 times as slow.

3 Gibbs-Markov Random Fields

If the graph is fixed and the node variable is binary, then we can write an ising model as


$$\text{net} \sim \text{nodeMatch}(\text{"var"}) + \text{nodeCount}(\text{"var"}) \mid \text{noDyad} + \text{var}$$

```
> form <- latticeNet ~ nodeMatch("var") + nodeCount("var") | noDyad + var
> n <- 25
> nw <- network.initialize(n, directed = FALSE)
> #create a toroidal lattice
> latticeNet <- as.BinaryNet(nw)
> latticeNet[["var"]] <- sample(c("a", "b"), n, replace=TRUE)
> nt <- function(i, j) sqrt(n)*(i-1) + (j-1) + 1
> for(node in 1:n){
+   i <- (node-1) %% sqrt(n) + 1
```

```

+   j <- floor((node-1)/sqrt(n)) + 1
+   latticeNet[node,nt(i,((j-2) %% sqrt(n)) + 1)] <- TRUE
+   latticeNet[node,nt(i,((j) %% sqrt(n)) + 1)] <- TRUE
+   latticeNet[nt((i-2) %% sqrt(n)) + 1 , j],node] <- TRUE
+   latticeNet[nt(((i) %% sqrt(n)) + 1 , j),node] <- TRUE
+ }
> plot(latticeNet)

```



Here we will simulate a model right on the boarder of the phase transition.
and fit a Gibbs field model to a point near the phase transition

```

> sampler <- createCppSampler(form,nodeSamplingPercentage=1)
> sampler$getModel()$setThetas(c(.9,0))
> samp <- sampler$generateSampleStatistics(1000,1000,10000)
> par(mfrow=c(1,2))
> hist(samp[,1],main="nodeMatch")
> hist(samp[,2],breaks=seq(from=-.5,to=25.5,length.out=27),main="nodeCount")
> colMeans(samp)

nodeMatch nodeCount
  44.6744   12.2569

> #fit the ernm
> fit <- ernm(form,meanStats=c(44.6216, 12.6480),mcmcInterval=1000)

```

```

sample statistics:
  means:
    nodeMatch nodeCount
simulated    25.0368    12.4504
observed     44.6216    12.6480
  std:
    nodeMatch nodeCount
[1,]    3.52759    2.514632

```

```

iteration: 1
log likelihood improved by: 9.351325
maximum scaled gradient: 5.551893
nodeMatch nodeCount
0.58806027 0.02288115
sample statistics:
  means:
    nodeMatch nodeCount
simulated    34.5662    13.3854
observed     44.6216    12.6480
  std:
    nodeMatch nodeCount
[1,]    5.1472    6.086557

```

```

iteration: 2
log likelihood improved by: 1.530644
maximum scaled gradient: 1.953567
nodeMatch nodeCount
0.906488944 -0.002949878
sample statistics:
  means:
    nodeMatch nodeCount
simulated    44.7414    12.3713
observed     44.6216    12.6480
  std:
    nodeMatch nodeCount
[1,]    5.146394    10.76868

```

```

iteration: 3
log likelihood improved by: 0.0005988971
maximum scaled gradient: 0.02569489

```



```

      nodeMatch      nodeCount
0.9020091088 -0.0005553294
sample statistics:
  means:
      nodeMatch nodeCount
simulated   44.6294    12.253
observed    44.6216    12.648
  std:
      nodeMatch nodeCount
[1,]  5.227293  10.71441

```

```

iteration: 4
log likelihood improved by:  0.0006802448
maximum scaled gradient:  0.03686622

```

```
> fit
```

ERNM Model

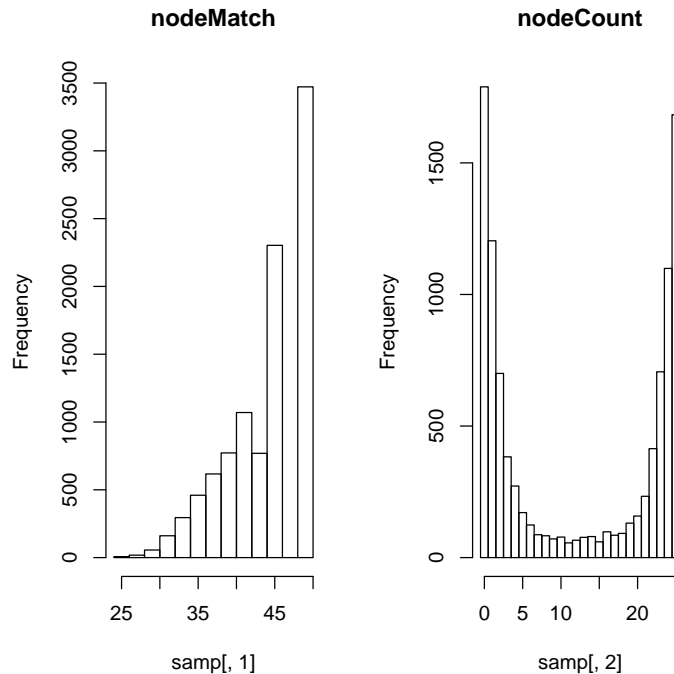
Domain:

Random graph = FALSE, Random variables = var

```

      nodeMatch      nodeCount
Parameters  0.9020091 -0.0005553294
Mean Values 44.6294000 12.2530000000

```



4 Working with BinaryNets

BinaryNets are network data structures native to the `ernm` package. They are special in a couple ways. First, they have a sparse representation of missingness, such that a directed network where half of the nodes have been egocentrically sampled takes the same amount of space as a simple fully observed network. Secondly, they may be passed up and down easily from R to C++ and visa versa. Finally, they are extensible on the C++ level, so that their implementation may be replaced. For example, it might be useful to put in a file backed storage system for large problems.

BinaryNets can be coerced to and from network objects.

```
> data(sampson)
> #coersion
> net <- as.BinaryNet(samplike)
> nw2 <- as.network(net)
> print(nw2)
```

```
Network attributes:
vertices = 18
directed = TRUE
hyper = FALSE
```

```

loops = FALSE
multiple = FALSE
bipartite = FALSE
total edges= 88
  missing edges= 0
  non-missing edges= 88

Vertex attribute names:
  cloisterville group vertex.names

> #dyad Extraction
> net[1:2,1:5]

      [,1] [,2] [,3] [,4] [,5]
[1,] FALSE TRUE TRUE FALSE TRUE
[2,] FALSE FALSE TRUE FALSE TRUE

> net$outNeighbors(c(1,2,3))

[[1]]
[1]  2  3  5  7 11 15

[[2]]
[1]  3  5  6  9 15

[[3]]
[1]  2  7  8 14

> #dyad assignment
> net[1,1:5] <- rep(NA,5)
> net[1:2,1:5]

      [,1] [,2] [,3] [,4] [,5]
[1,] FALSE  NA  NA  NA  NA
[2,] FALSE FALSE TRUE FALSE TRUE

> net[1:2,1:5,maskMissing=FALSE] #remove the mask over missing values and see nothing was re

      [,1] [,2] [,3] [,4] [,5]
[1,] FALSE TRUE TRUE FALSE TRUE
[2,] FALSE FALSE TRUE FALSE TRUE

> #node variables
> net$variableNames()

$discrete
[1] "group"          "na"              "vertex.names"

$continuous
[1] "cloisterville"

```

```

> net[["group"]]

[1] Loyal    Loyal    Loyal    Loyal    Loyal    Loyal    Loyal    Turks
[9] Turks    Turks    Turks    Turks    Turks    Turks    Outcasts Outcasts
[17] Outcasts Outcasts
Levels: Loyal Outcasts Turks

> net[["rnorm"]] <- rnorm(18)
> net[["rnorm"]]

[1] 0.14441451 -0.57949376 -3.47592196 -0.94085147 0.37996743 -0.11105589
[7] -1.14843258 0.59404012 0.56828957 -1.20534168 0.57650561 -0.78246131
[13] 1.19329181 0.86643758 -1.08681315 -0.05110687 -0.74386960 -0.57519834

>
> #See available methods
> #print(DirectedNet)
> #print(UndirectedNet)

```

5 A simple ERNM for Sampson's Monks

Here we will create a model with both the graph and group being random.

```

> net <- as.BinaryNet(samplike)
> form <- net ~ edges() + reciprocity() + homophily("group") + nodeCount("group") | group
> fit <- ernm(form,verbose=0)

```

calculating:

```
> fit
```

ERNM Model

Domain:

Random graph = TRUE, Random variables = group


	edges	reciprocity	homophily	nodeCount1	nodeCount2
Parameters	-1.148926	1.365997	5.667398	0.02746373	-2.222304
Mean Values	88.208500	28.064900	15.311497	7.09270000	3.992500

```


> #generate some networks from the model
> sampler <- createCppSampler(form)
> sampler$getModel()$setThetas(fit$theta)
> netList <- sampler$generateSample(1000,1000,3)
> par(mfrow=c(2,2))
> plot(net,main="Sampson's Monks")
> plot(netList[[1]],main="sim 1")
> plot(netList[[2]],main="sim 1")
> plot(netList[[3]],main="sim 1")

```


Sampson's Monks




sim 1



sim 1



sim 1



6 A latent class model for Sampson's Monks

```
> net <- as.BinaryNet(samplike)
> net[["latent"]] <- factor(rep(NA,18),levels=c("a","b","c"))
> form <- net ~ edges() + reciprocity() + homophily("latent") + nodeCount("latent") | latent
> fit <- ernm(form,modelType=MissingErnmModel,verbose=0,theta0=c(-1.8,2.4,1,0,0))
```

calculating:

```
> fit
```

ERNM Model (with missing values)

Domain:

Random graph = TRUE, Random variables = latent

	edges	reciprocity	homophily	nodeCount1	nodeCount2
Parameters	-1.125041	1.323231	5.773252	0.2307555	-2.216451
Mean Values	88.607600	28.107400	15.414312	7.2210000	4.008200

```
> #simulate 5 networks conditional upon observed graph
> sampler <- createCppSampler(form,sampler="NTDNBRNonObservedMetropolis"
+                               ,nodeSamplingPercentage=1)
```

```

> sampler$getModel()$setThetas(fit$theta)
> netList <- sampler$generateSample(1000,1000,5)
> #all simulations yield identical clusters
> cbind(SampsonsClusters=as.character(net[["group"]]),
+       Simulations=sapply(netList,function(x) x$getVariable("latent",maskMissing=FALSE)))

```

```

      SampsonsClusters
[1,] "Loyal"          "b" "b" "b" "b" "b"
[2,] "Loyal"          "b" "b" "b" "b" "b"
[3,] "Loyal"          "b" "b" "b" "b" "b"
[4,] "Loyal"          "b" "b" "b" "b" "b"
[5,] "Loyal"          "b" "b" "b" "b" "b"
[6,] "Loyal"          "b" "b" "b" "b" "b"
[7,] "Loyal"          "b" "b" "b" "b" "b"
[8,] "Turks"          "a" "a" "a" "a" "a"
[9,] "Turks"          "a" "a" "a" "a" "a"
[10,] "Turks"         "a" "a" "a" "a" "a"
[11,] "Turks"         "a" "a" "a" "a" "a"
[12,] "Turks"         "a" "a" "a" "a" "a"
[13,] "Turks"         "a" "a" "a" "a" "a"
[14,] "Turks"         "a" "a" "a" "a" "a"
[15,] "Outcasts"      "c" "c" "c" "c" "c"
[16,] "Outcasts"      "c" "c" "c" "c" "c"
[17,] "Outcasts"      "c" "c" "c" "c" "c"
[18,] "Outcasts"      "c" "c" "c" "c" "c"

```